

User's Guide to Running the Draft NIST SP 800-90B Section 9 Entropy Estimation Tests

30 September 2014

T. A. Hall

This is a brief introduction on how to run the Python command-line programs that implement the statistical entropy estimation tests found in Section 9 of the Draft NIST SP 800-90B (August 2012). It is not a description or explanation of the tests themselves. Please refer to the standard itself for definitions and descriptions of the tests and their rationales.

Installation

Uncompress the file *py90b_v4.zip* into a single directory. The contents of the zip archive are:

Python files to implement the tests:

iidmain.py

- Contains main routine to give the IID entropy estimate, if IID assumption holds
- Run six shuffle tests to determine if IID
- Run chi-square independence and goodness of fit tests to determine if IID
- Estimate min entropy if passes above tests
- Run sanity check tests

noniid_main.py

- Contains main routine to give the non-IID entropy estimate
- Run five tests to estimate min-entropy
- min-entropy as lowest of the five
- Run sanity check tests

shuffle_tests.py

- Contains six shuffle tests to determine if dataset is IID from Section 9.1.2

chi_square_tests.py

- Contains the chi square independence and goodness of fit for binary and non-binary data from Section 9.1.3

iid_tests.py

- Contains the min entropy calculation for IID data from Section 9.2

sanity_checks.py

- Contains the two sanity check tests from Section 9.4

noniid_collision.py

- the non-IID collision test from Section 9.3.3
- Test may not be valid for all datasets.

partial_collection.py

- The non-IID partial collection test from Section 9.3.4
- Test may not be valid for all datasets.

markov.py

- Contains the non-IID Markov test from Section 9.3.5
- Per SP 800-90B, only up to 6 bits per symbol used for Markov test

maurer.py

- Contains the non-IID compression test (Maurer Universal Statistic) from Section 9.3.6
- Test may not be valid for all datasets

frequency.py

- Contains the non-IID frequency test from Section 9.3.7

util90b.py

- Contains utility functions

Sample dataset files:

Three files generated with TrueRand that should pass determine iid tests:

- 1 000 000 data samples
 - 1 bit per sample (***truerand_1bit.bin***)
 - 4 bits per sample (***truerand_4bit.bin***)
 - 8 bits per sample (***truerand_8bit.bin***)

One file generated with TrueRand that should pass shuffle tests but fail chi square tests:

- 1 000 000 data samples, 9 bits per sample (***truerand_9bit.bin***)

One file containing binary digits of pi that fails iid tests:

- ***data.pi.bin***

This User Guide

- *user_guide.pdf*

The code has been tested and run successfully on the following:

- Python 2.6 on Linux
- Python 2.7 on Mac OS X
- Python 3.3 on Windows 7

It should run on any OS with Python 2.6+ or Python 3.

Note that this tool does not come with a Python installation. If you do not already have Python installed on your system, go to <https://www.python.org> and select “Download.” No additional modules or packages are required to run the code. However, some routines will run faster if you have the **numpy** package installed. You can get **numpy** at <http://www.scipy.org>. If you are running a Windows OS, you can also find it here: <http://www.lfd.uci.edu/~gohlke/pythonlibs>. Alternatively, you can download the entire **scipy-stack**, which includes **numpy**.

Running the tests

The help message for the IID tests is:

```
C:\est>python iidmain.py -h
usage: iid_main.py [-h] [-v] datafile bits_per_symbol number_of_shuffles
Run the Draft NIST SP 800-90B (August 2012) IID Tests

positional arguments:
  datafile              dataset on which to run tests
  bits_per_symbol       number of bits used to represent sample output values
  number_of_shuffles    number of shuffles per data subset for shuffle tests

optional arguments:
  -h, --help            show this help message and exit
  -v, --verbose          verbose mode: show detailed test results
```

Examples of running the IID tests follow.

Run the IID tests on the included truerand_8bit.bin dataset, which contains 8 bits per sample. Use 1,000 shuffles of the data subsets and append the verbose flag for detailed output:

```
C:\est>python iid_main.py truerand_8bit.bin 8 1000 -v
Read in file truerand_8bit.bin, 1000000 bytes long.
Dataset: 1000000 8-bit symbols.
Output symbol values: min = 0, max = 255
```

Compression Test:

Scores	Ranks
106842	670
106886	856
106858	673
106718	106
106845	564
106899	867
106752	224
106867	741
106936	981*
106849	695

	1

Passed Compression Test

...

The full program output is not listed for space considerations. The first three lines of output are information about the dataset: its name, total size in bytes, how the raw bytes are interpreted (1,000,000 8-bit symbols as opposed to 500,000 16-bit symbols, for example) and the range of sample values in the dataset.

Following this is detailed information about the individual shuffle tests. The test name is followed by the scores and ranks of the 10 original (unshuffled) data subsets. If the rank of the score an original (unshuffled) data subset in a ranked ordering of the scores of all 1,000 shuffled data subsets is in the top or bottom 5%, then the rank is marked with an asterisk. For 1,000 shuffles, this works out to ranks of greater than or equal 950 or less than or equal to 50. If eight or more of the data subsets fall in this range, then the test fails, indicating that the IID assumption does not hold. Please see the Draft NIST SP 800-90B (August 2012) for an explanation and more details on this. A similar display of scores, ranks and Passed/Failed verdict is output for the other five shuffle tests.

If the dataset passes all of the shuffle tests, as is the case for truerand_8bit.bin, then the program output indicates this and prints out details and results of the Chi-square tests and the overall determination of the IID assumption. If the determination is that the IID assumption holds, as is true for our example, the min-entropy estimate is output, followed by the details and results of the two sanity checks.

```
** Passed iid shuffle tests
Chi square independence
    score = 65212.5, degrees of freedom = 65280, cut-off = 66402.2
** Passed chi-square independence test

Chi square stability
    score = 2449.48, degrees of freedom = 2313 cut-off = 2528.88
** Passed chi-square stability test

IID = True
min-entropy = 7.87108

Compression sanity check...
    dataset 1 compressed length = 854736, cutoff = 787108...Pass

    dataset 2 compressed length = 855088, cutoff = 787108...Pass

    dataset 3 compressed length = 854864, cutoff = 787108...Pass

    dataset 4 compressed length = 853744, cutoff = 787108...Pass

    dataset 5 compressed length = 854760, cutoff = 787108...Pass

    dataset 6 compressed length = 855192, cutoff = 787108...Pass

    dataset 7 compressed length = 854016, cutoff = 787108...Pass

    dataset 8 compressed length = 854936, cutoff = 787108...Pass

    dataset 9 compressed length = 855488, cutoff = 787108...Pass

    dataset 10 compressed length = 854792, cutoff = 787108...Pass

Collision sanity check...
    Dividing dataset into 4-tuples
    Check rule 1 - do three or more 4-tuples have the same value?...Pass
    Check rule 2 - probability of number of collisions below cutoff
        number of collisions = 6, cutoff = 10.4023...Pass

sanity check = PASS
```

If the same dataset and test parameters are run without the verbose flag, only the results of the IID determination, min-entropy estimate and sanity check are output:

```
C:\est>python iid_main.py truerand_8bit.bin 8 1000
IID = True
min-entropy = 7.87108
sanity check = PASS
```

Note that 1,000 is used as the number of shuffles for the examples above. 1,000 shuffles is specified in Draft NIST SP 800-90B (August 2012) and thus what should be used in order to run the tests in conformance with the standard. However, you may use a different number if you choose.

The help message for the non-IID tests is:

```
C:\est>python noniid_main.py -h
usage: noniid_main.py [-h] [-u use_bits] [-v] datafile bits_per_symbol

Run the Draft NIST SP 800-90B (August 2012) non-IID Tests

positional arguments:
  datafile              dataset on which to run tests
  bits_per_symbol       number of bits used to represent sample output values

optional arguments:
  -h, --help            show this help message and exit
  -u use_bits, --usebits use_bits
                        use only the N lowest order bits per sample
  -v, --verbose         verbose mode: show detailed test results
```

Next are some examples of running the non-IID tests.

First example: run the non-IID tests on the included truerand_4bit.bin dataset, which has 4 bits per sample. The verbose flag is set in order to obtain detailed test results. Note that the results of each individual non-IID test are shown. The output is below:

```

C:\est>python noniid_main.py truerand_4bit.bin 4 -v
Read in file truerand_4bit.bin, 1000000 bytes long.
Dataset: 1000000 4-bit symbols.
Output symbol values: min = 0, max = 15
- Collision test          : p(max) = 0.0715332, min-entropy = 3.80524
- Partial collection test : p(max) = 0.074295, min-entropy = 3.75059
- Markov test            : p(max) = 3.02369e-153, min-entropy = 3.95827
- Compression test       : p(max) = 0.0789795, min-entropy = 3.66238
- Frequency test         : p(max) = 0.063134, min-entropy = 3.95774
min-entropy = 3.66238

Compression sanity check...
  dataset 1 compressed length = 435232, cutoff = 366238...Pass

  dataset 2 compressed length = 435720, cutoff = 366238...Pass

  dataset 3 compressed length = 435904, cutoff = 366238...Pass

  dataset 4 compressed length = 435336, cutoff = 366238...Pass

  dataset 5 compressed length = 435768, cutoff = 366238...Pass

  dataset 6 compressed length = 435480, cutoff = 366238...Pass

  dataset 7 compressed length = 435288, cutoff = 366238...Pass

  dataset 8 compressed length = 435632, cutoff = 366238...Pass

  dataset 9 compressed length = 435648, cutoff = 366238...Pass

  dataset 10 compressed length = 435936, cutoff = 366238...Pass

Collision sanity check...
  Dividing dataset into 9-tuples
  Check rule 1 - do three or more 9-tuples have the same value?...Pass
  Check rule 2 - probability of number of collisions below cutoff
    number of collisions = 0, cutoff = 0.738097...Pass

sanity check = PASS

```

As with the IID tests, not setting the verbose flag produces compact results output:

```
C:\est>python noniid_main.py truerand_4bit.bin 4
min-entropy = 3.66238
sanity check = PASS
```

The usebits option allows you to instruct the program to consider only a lower order subset of the bits per sample. This is useful when almost all of the entropy is in these low order bits. Below is the case where only the lowest order two bits of the four bit samples are used.

```
C:\est>python noniid_main.py truerand_4bit.bin 4 --usebits 2 -v
Read in file truerand_4bit.bin, 1000000 bytes long.
Dataset: 1000000 4-bit symbols.
Output symbol values: min = 0, max = 15
* Using only low 2 bits out of 4.
* Using output symbol values: min = 0, max = 3
- Collision test          : p(max) = 0.277344, min-entropy = 1.85025
- Partial collection test : p(max) = 0.276367, min-entropy = 1.85534
- Markov test            : p(max) = 1.48356e-77, min-entropy = 1.9939
- Compression test       : p(max) = 0.276001, min-entropy = 1.85725
- Frequency test         : p(max) = 0.250906, min-entropy = 1.98776
min-entropy = 1.85025

Compression sanity check...
    dataset 1 compressed length = 215816, cutoff = 185025...Pass

    ...<output deleted to save space>...

    dataset 10 compressed length = 215928, cutoff = 185025...Pass

Collision sanity check...
    Dividing dataset into 17-tuples
    Check rule 1 - do three or more 17-tuples have the same value?...Pass
    Check rule 2 - probability of number of collisions below cutoff
        number of collisions = 0, cutoff = 0.587999...Pass

sanity check = PASS
```

Disclaimers

- This code is made available without any assertion or guarantee, implied or otherwise, of correctness or completeness.

- No support is provided for this code.
- The identification of any commercial product or trade name does not imply endorsement or recommendation by the National Institute of Standards and Technology, nor is it intended to imply that the materials or equipment identified are necessarily the best available for the purpose.

Appendix: Selected Runs of IID and Non-IID Tests

Non-IID Tests on 1 bit TrueRand data:

```
C:\est>python noniid_main.py truerand_1bit.bin 1 -v
Read in file truerand_1bit.bin, 1000000 bytes long.
Dataset: 1000000 1-bit symbols.
Output symbol values: min = 0, max = 1
- Collision test          : p(max) = 0.53125, min-entropy = 0.912537
- Partial collection test : p(max) = 0.526367, min-entropy = 0.925859
- Markov test            : p(max) = 3.39255e-39, min-entropy = 0.998381
- Compression test       : p(max) = 0.529541, min-entropy = 0.917186
- Frequency test         : p(max) = 0.500433, min-entropy = 0.995227
min-entropy = 0.912537

Compression sanity check...
  dataset 1 compressed length = 125008, cutoff = 91253.7...Pass
  dataset 2 compressed length = 124792, cutoff = 91253.7...Pass
  dataset 3 compressed length = 124920, cutoff = 91253.7...Pass
  dataset 4 compressed length = 124896, cutoff = 91253.7...Pass
  dataset 5 compressed length = 124904, cutoff = 91253.7...Pass
  dataset 6 compressed length = 124824, cutoff = 91253.7...Pass
  dataset 7 compressed length = 124912, cutoff = 91253.7...Pass
  dataset 8 compressed length = 124920, cutoff = 91253.7...Pass
  dataset 9 compressed length = 125032, cutoff = 91253.7...Pass
  dataset 10 compressed length = 125240, cutoff = 91253.7...Pass

Collision sanity check...
  Dividing dataset into 32-tuples
  Check rule 1 - do three or more 32-tuples have the same value?...Pass
  Check rule 2 - probability of number of collisions below cutoff
    number of collisions = 0, cutoff = 0.791109...Pass

sanity check = PASS
```

Non-IID Tests on 9-bit TrueRand data:

```
C:\est>python noniid_main.py truerand_9bit.bin 9 -v
Read in file truerand_9bit.bin, 2000000 bytes long.
Dataset: 1000000 9-bit symbols.
Output symbol values: min = 0, max = 511
- Collision test          : p(max) = 0.00601578, min-entropy = 7.37703
- Partial collection test : p(max) = 0.00437081, min-entropy = 7.83788
- Markov test (map 6 bits): p(max) = 5.70333e-223, min-entropy = 5.7678
- Compression test       : p(max) = 0.00553894, min-entropy = 7.49617
- Frequency test         : p(max) = 0.002091, min-entropy = 8.23683
min-entropy = 5.7678
```

Compression sanity check...

```
dataset 1 compressed length = 948384, cutoff = 576780...Pass

dataset 2 compressed length = 947576, cutoff = 576780...Pass

dataset 3 compressed length = 946792, cutoff = 576780...Pass

dataset 4 compressed length = 949304, cutoff = 576780...Pass

dataset 5 compressed length = 947480, cutoff = 576780...Pass

dataset 6 compressed length = 947896, cutoff = 576780...Pass

dataset 7 compressed length = 950720, cutoff = 576780...Pass

dataset 8 compressed length = 947952, cutoff = 576780...Pass

dataset 9 compressed length = 946872, cutoff = 576780...Pass

dataset 10 compressed length = 947384, cutoff = 576780...Pass
```

Collision sanity check...

```
Dividing dataset into 6-tuples
Check rule 1 - do three or more 6-tuples have the same value?...Pass
Check rule 2 - probability of number of collisions below cutoff
number of collisions = 0, cutoff = 0.530863...Pass
```

sanity check = PASS

IID Tests run on binary digits of pi

Note that since the reordering of samples in the shuffle tests is random, the ranks may change from run to run of the IID tests for the same dataset.

```
C:\est>python iid_main.py data.pi.bin 1 1000 -v
```

```
Read in file data.pi.bin, 1165666 bytes long.
```

```
Dataset: 1165666 1-bit symbols.
```

```
Output symbol values: min = 0, max = 1
```

Compression Test:

Scores	Ranks
17867	1*
17865	1*
17874	1*
17904	1*
17851	1*
17846	1*
17894	1*
17888	1*
17877	1*
17910	1*

	10

Failed Compression Test

Over/under Test:

Scores	Ranks
22 52034	826 1*
19 52410	501 1*
19 52318	501 1*
20 52577	549 1*
20 52347	501 1*
18 52090	328 1*
20 52132	547 1*
19 52169	501 1*
19 52136	501 1*
19 52131	501 1*

	0 10

Failed Over/under Test

Excursion Test:

Scores	Ranks
156.625	642
125.619	364
160.774	699
138.821	507
84.1666	42*
85.0692	54
114.445	229
93.5417	80
102.259	123
109.691	238

	1

Passed Excursion Test

Directional runs Test:

Scores			Ranks		
8735	8	6102	1000*	187	1000*
8721	8	6043	1000*	202	1000*
8710	9	6090	1000*	501	1000*
8734	9	6072	1000*	501	1000*
8779	8	6096	1000*	182	1000*
8810	9	6088	1000*	501	1000*
8711	8	6046	1000*	162	1000*
8736	8	6064	1000*	190	1000*
8842	9	6158	1000*	501	1000*
8636	9	6086	979*	501	1000*
			---	---	---
			10	0	10

Failed Directional runs Test

Covariance Test:

Scores	Ranks
0.0219853	1000*
0.0203951	1000*
0.0210892	1000*
0.0200727	1000*
0.0205968	1000*
0.021612	1000*
0.0218409	1000*
0.0212882	1000*
0.0217061	1000*
0.0216055	1000*

	10

Failed Covariance Test

Collision Test:

Scores			Ranks		
1	18.4866	56	501	626	635
1	18.9863	49	501	964*	124
1	18.051	46	501	146	13*
1	17.9182	56	501	59	622
1	18.2407	49	501	390	132
1	17.9208	47	501	116	38*
1	18.5787	52	501	662	370
1	17.9169	51	501	124	276
2	17.9608	56	971*	102	634
1	18.0984	49	501	235	135
---			---	---	---
			1	1	2

Passed Collision Test

** Failed iid shuffle tests

IID = False